Lesson 14: Idea Analytics

Managing Innovation - Master in Innovation Management and Business Development - Spring 2021

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Center for Humanities Computing
Aarhus University, Denmark

April 19, 2021

CENTER FOR HUMANITIES
COMPUTING AARHUS

Outline

1. Data*
   - Data
   - Data science
   - KDD
   - Normalization
   - Big Data?
   - Human-in-the-Loop Models
2. Example
   - Trend Reservoirs
3. Techniques
   - Text analytics
   - Word counts
   - Word distributions
   - Vector space model
   - Sentiment analysis
   - Non-negative Matrix Factorization
   - Classification
   - EDA
   - Neural embeddings

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Data Access

`https://worker02.chcaa.au.dk/jupyter/hub/login`

Communication

`chcaa.cas.au.dk`

Add subject: [Group <your group number>] <description>

Example: You want to submit group 5's requirements:

- list the group's requirements
- add '[Group 5] requirements' to the subject line
- submit to `chcaa.cas.au.dk`

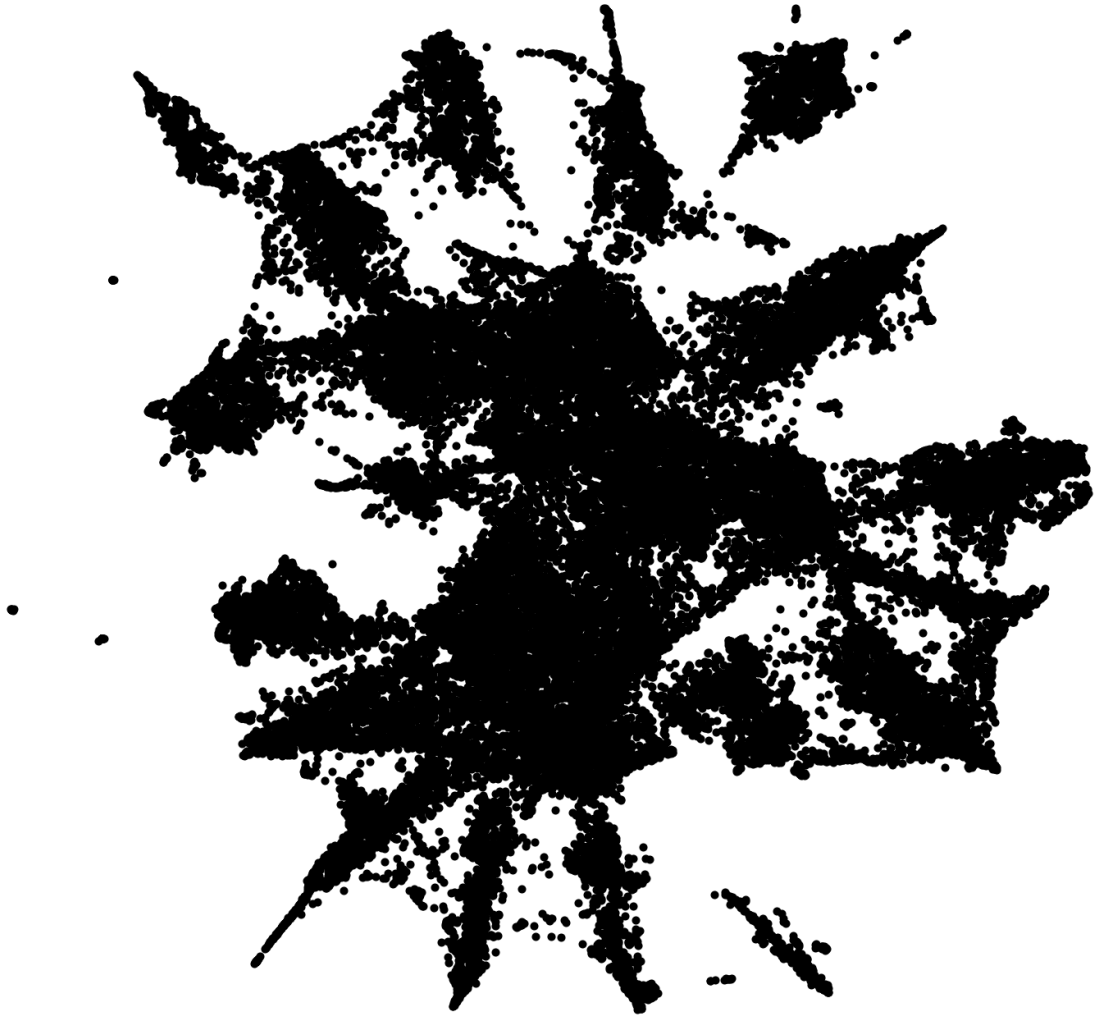

JSTOR Data For Research, $n = 43,802$, $\sim$ 400m words, Eng*, Fr, Ger.

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

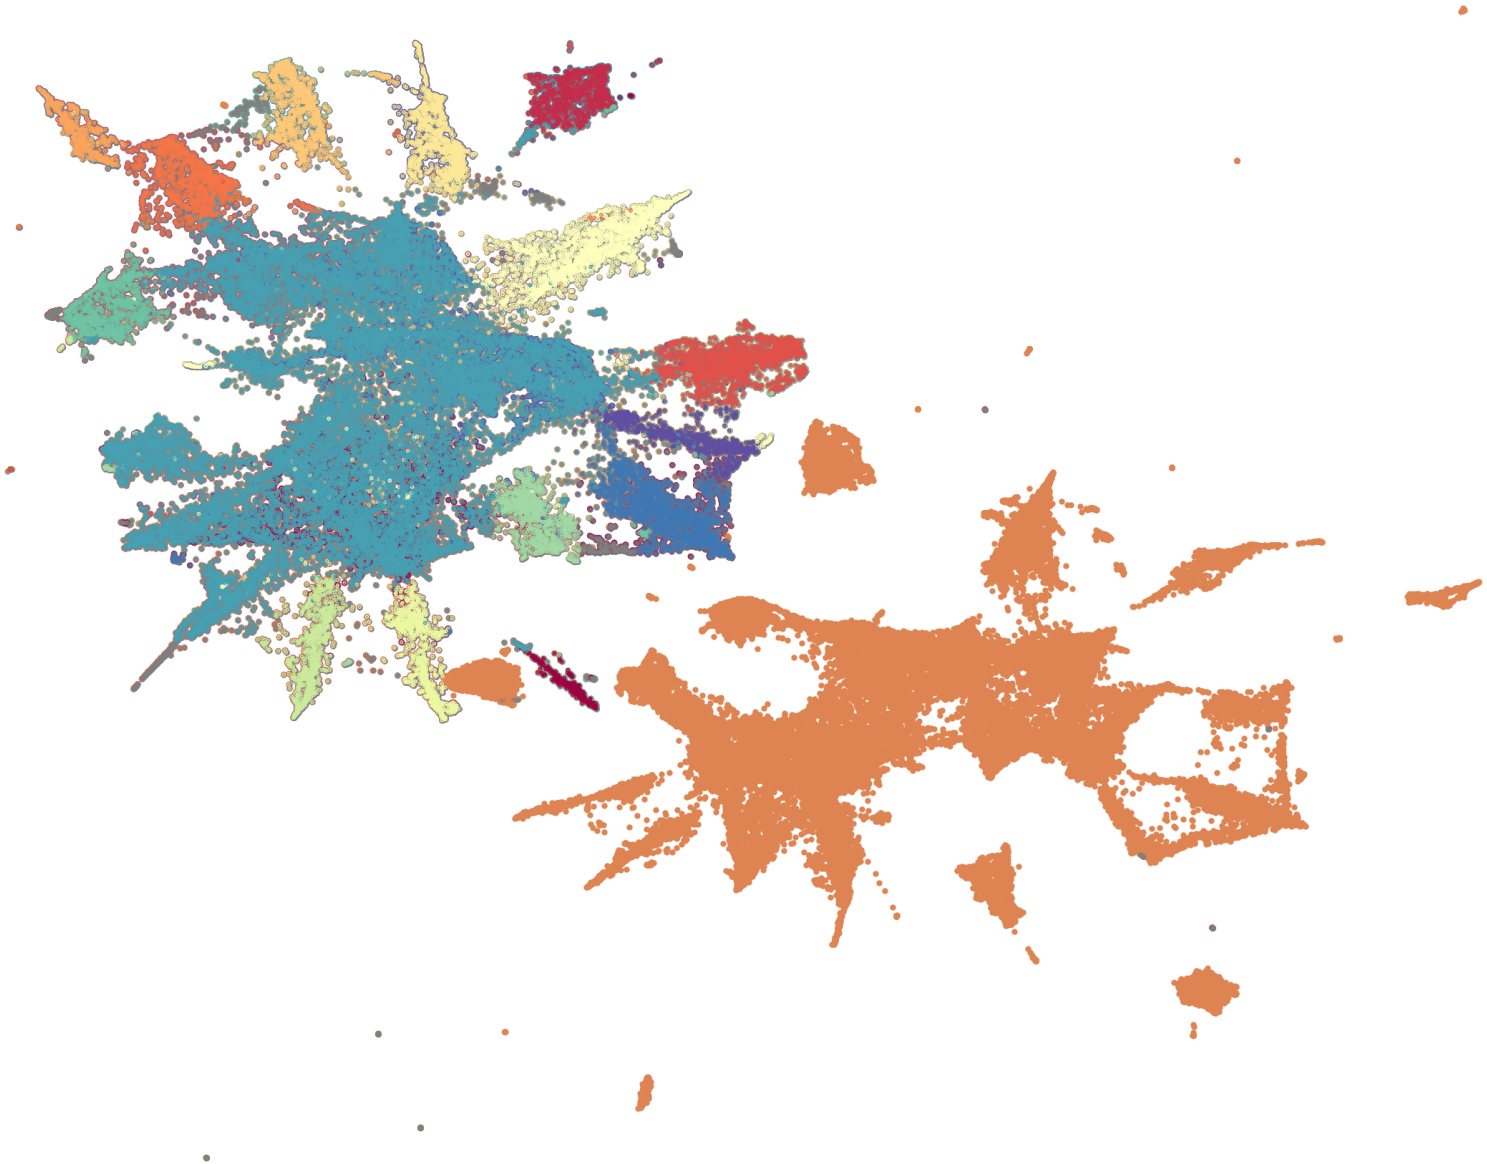

CENTER FOR HUMANITIES
COMPUTING AARHUS

Lesson 14: Idea Analytics

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Why does this matter?

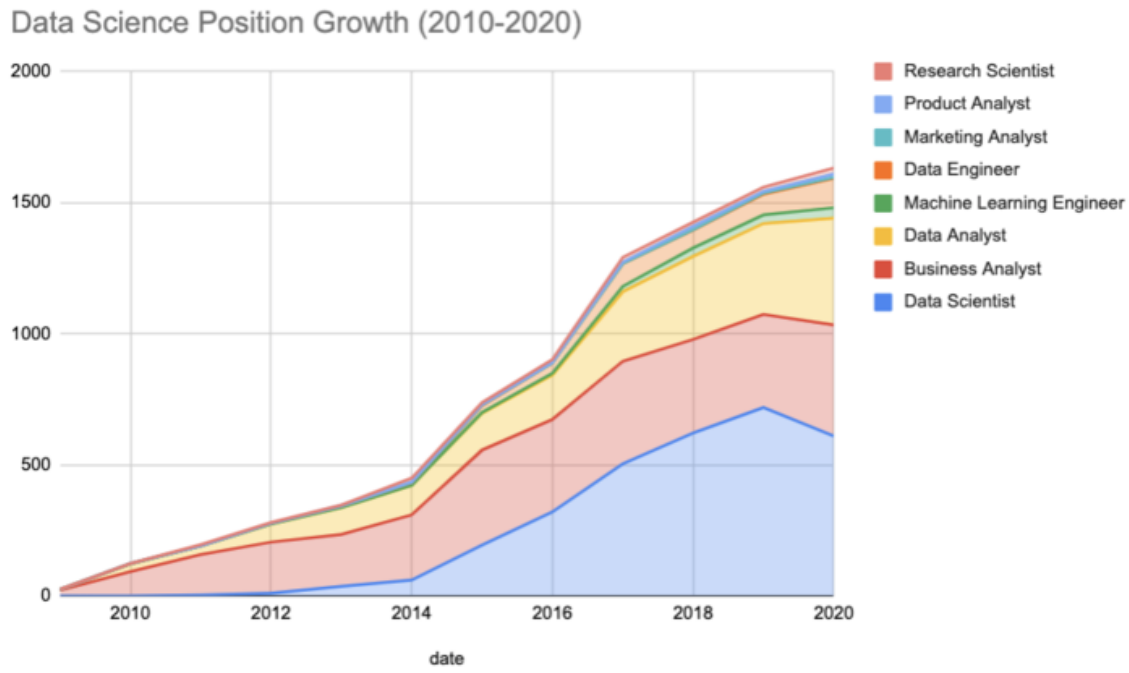

With the rise of Big Data, we have seen a rapid increase in data-related positions, source: Stephanie Glen (2021) 'Data Science Job Market Shrinking? Not So Fast'

Knowledge Discovery in Data(bases)

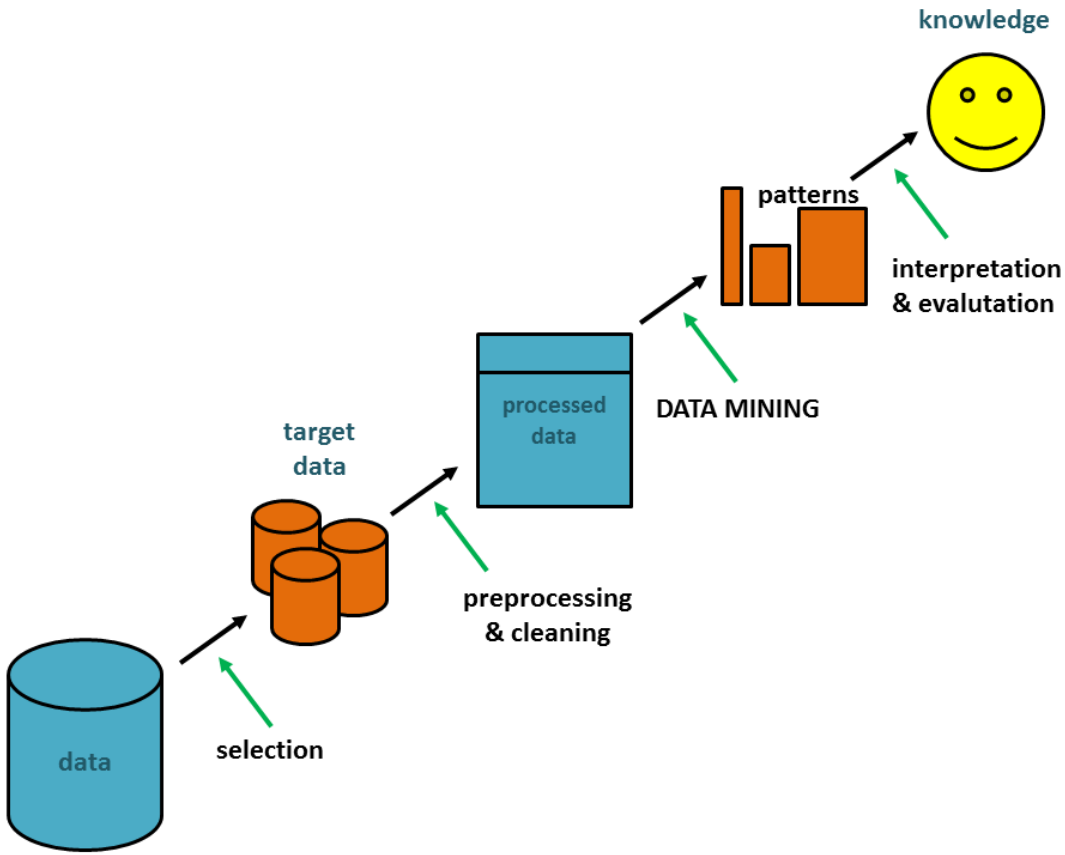

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Data Normalization

We normalize natural language data in order to increase our signal/reduce the randomness

- Removal of duplicate whitespaces and punctuation.
- Casefolding.
- Removal or substitution of special characters/emojis (e.g.: remove hashtags).
- Substitution of contractions (very common in English; e.g.: 'I'm'→'I am').
- Transform word numerals into numbers (eg.: 'twenty three'→'23').
- Substitution of values for their type (e.g.: '50DKK'→'MONEY').
- Acronym normalization (e.g.: 'DK'→'Denmark') and abbreviation normalization (e.g.: 'btw'→'by the way').
- Normalize date formats, social security numbers or other data that have a standard format.
- Spell correction
- Removal of gender/time/grade variation with Stemming or Lemmatization.
- Substitution of rare words for more common synonyms.
- Stop word removal (∼ dimensionality reduction technique).

Big Data are not necessarily enough data to solve your problem.

- wrong data
- lack precision/low detection rate
- problem complexity
- (too many data)

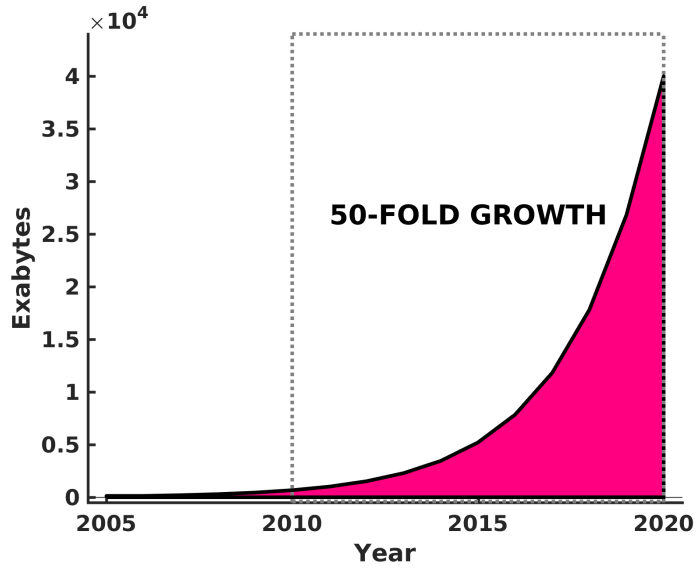

- revise question*
- collect more data
- enrich existing data*

Human-in-the-Loop Models

as task complexity increases, a need for (operational approaches to) leveraging human intelligence in the development of learning algorithms has become apparent

Type	Human Involvement	Resources	Relevance
Out-of-the-loop	not required	low	low
On-the-loop	checking	medium	medium↓
In-the-loop	required	high	medium↑

WHEN

algorithms are not understanding the input

data input is interpreted incorrectly

algorithms do not know how to perform the task

to make models more accurate

cost of errors is too high in development

data is rare or not available

THEN

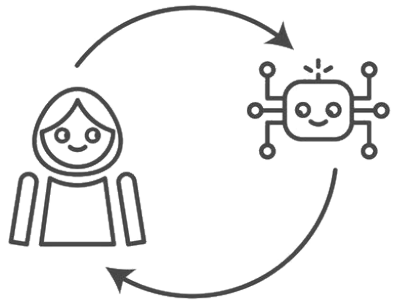

HITL Models

CENTER FOR HUMANITIES COMPUTING AARHUS

Trend Reservoirs in Social Media

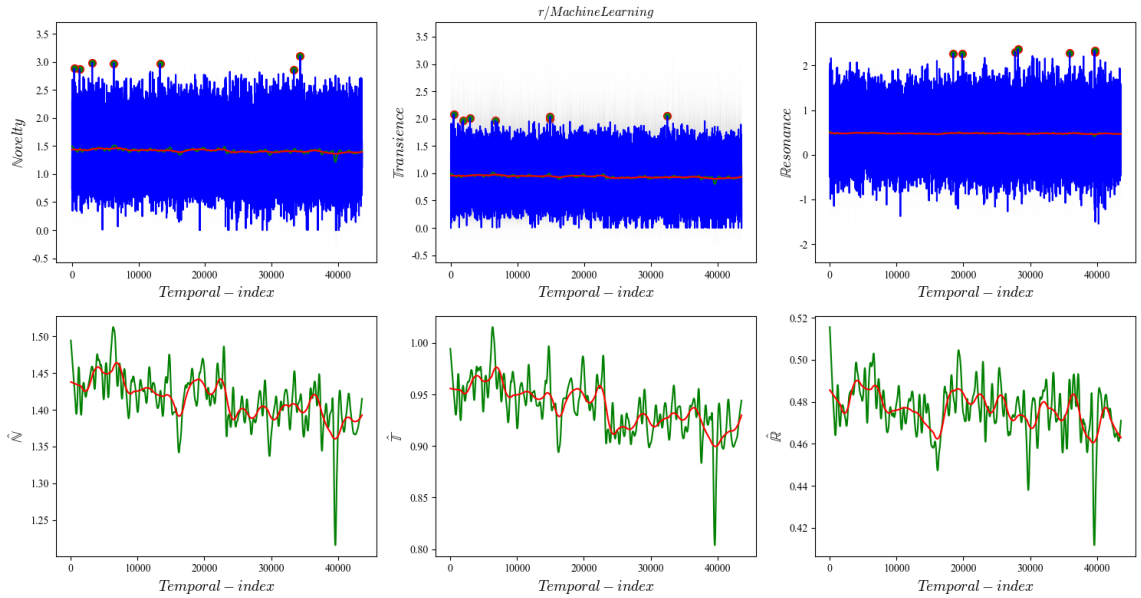

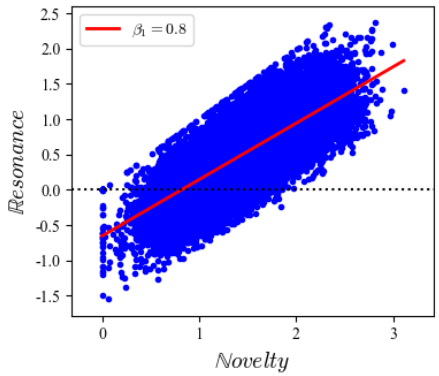

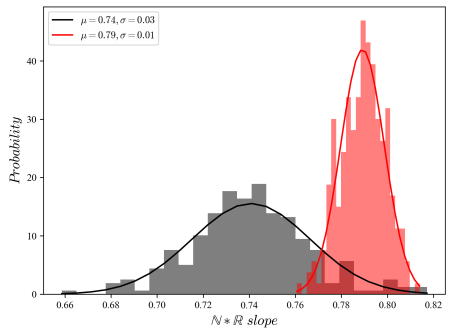

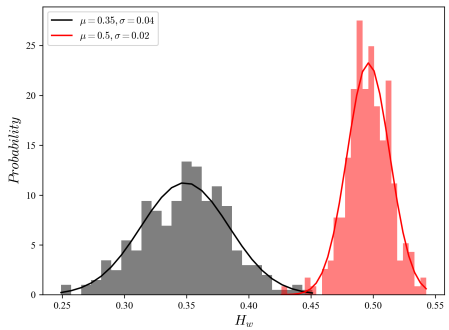

trend reservoirs (i.e., social media signals that display high trend potential) can be identified by their relationship between novel and resonant behavior, and their minimal persistence.

Text Analytics

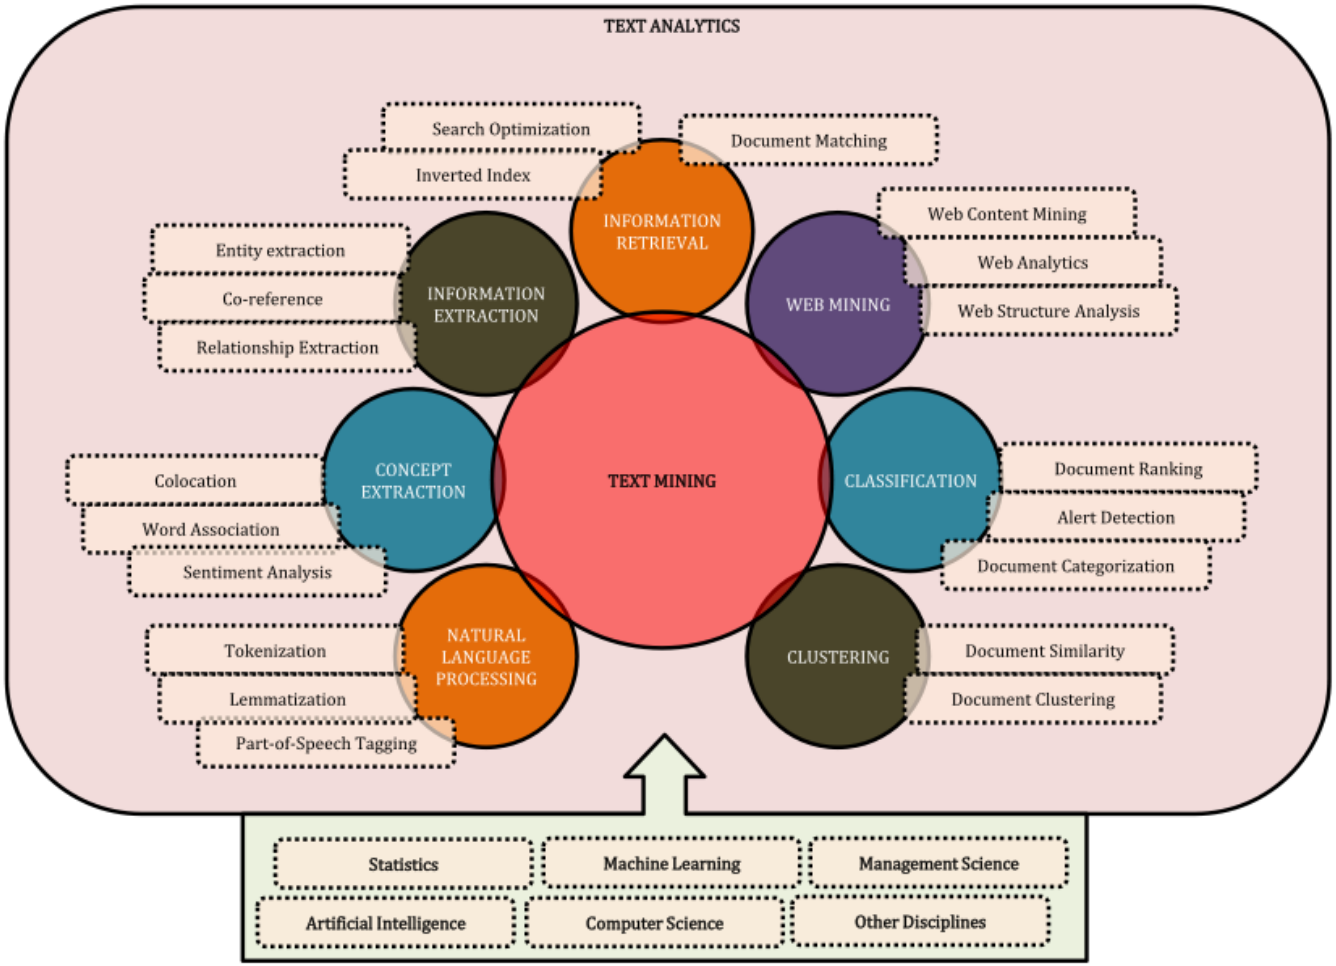

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Word counts

words are (one of) the basic units of meaning

most TM techniques rely on word frequencies, that is, we tokenize a text at the word level and count the number of tokens for each type

I am Daniel	
I am Sam	'I' 'am' 'Daniel' 'I'
Sam I am	'am' 'Sam' 'Sam' 'I'
That Sam-I-am	'am' 'That' 'Sam' 'I'
That Sam-I-am!	'am' 'That' 'Sam' 'I'
I do not like	'am' 'I' 'do' 'not' 'like'
that Sam-I-am	'that' 'Sam' 'I' 'am' ...
...	

a	1	59	0.073
am	1	16	0.02
and	1	24	0.03
anywhere	1	1	0.001
anywhere	1	7	0.009
...			
you	1	34	0.042
total	55	804	1.0

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Word Distributions

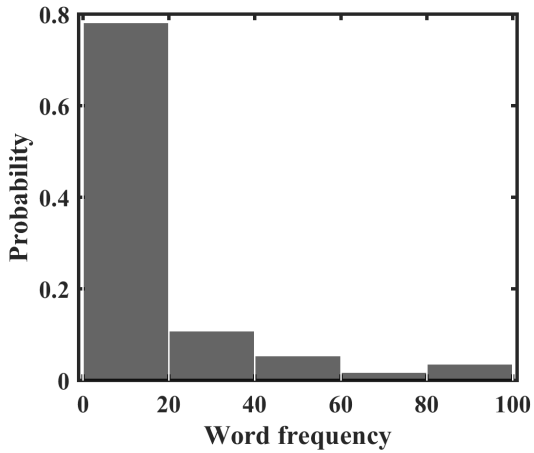

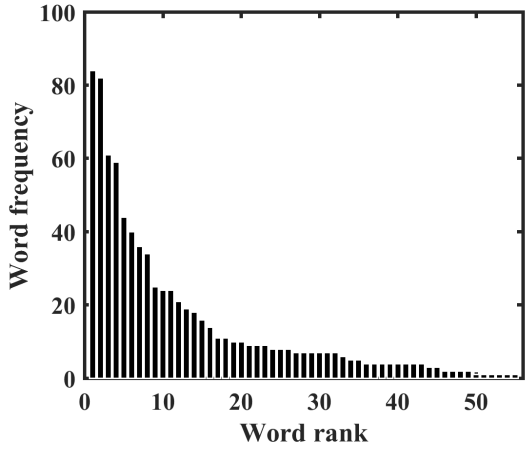

Most words are infrequent, but a few words are very frequent

'i' 'not' 'them' 'a' 'like' 'in' 'do' 'you' 'would'

Model a text as a distribution over words. Some words are more likely than other.

Often times we are looking at the mid-range (not too likely and not too unlikely).

Vector Space Model

any collection of m documents can be represented in the vector space model by a document-term matrix of m documents and n terms

a vector space model is a basic modeling mechanism for a word- or document-space (whether we look at rows or columns)

- a document vector with only one word is collinear to the vocabulary word axis
- a document vector that does not contain a specific word is orthogonal/perpendicular to the word axis
- two documents are identical if they contain the same words in a different order (BOW assumption)

You can use the 'words' as predictors $x_1, x_2, \ldots x_m$ in:

$$Y_i = f(X_i, \beta) + \epsilon_i$$

but it can result in an underdetermined system

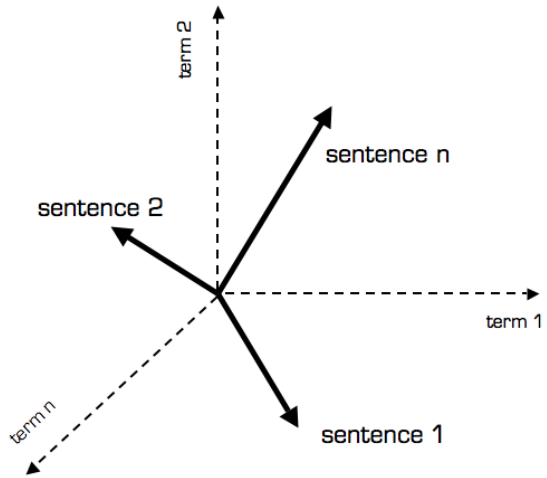

Document space	t_1	t_2	t_3	...	t_n ← Term vector space
D_1	a_{11}	a_{12}	a_{13}	...	a_{1n}
D_2	a_{21}	a_{22}	a_{23}	...	a_{2n}
D_3	a_{31}	a_{32}	a_{33}	...	a_{3n}
...					
D_m	a_{m1}	a_{m2}	a_{m3}	...	a_{mn}
Q	b_1	b_2	b_3	...	b_n

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Sentiment Analysis

- we can use a dictionary to extract cognitive and affective keywords from a collection of documents and apply a sentiment function

```
'Did Crooked Hillary help disgusting (check out sex tape and past) Alicia M become a U.S. citizen
so she could use her in the debate?'

Positive sex, citizen
Negative crooked, hillary, disgusting, out
Sentiment Score (2+1) + (-2-1-3-1) = -4
Sentiment Polarity Negative
Overall Score Sum of all sentence scores
```

- a sentiment vector is a vector of keyword frequencies weighted by sentiment scores

> You can avoid the underdetermined system and use the overall 'sentiment' as predictor in: $Y_i = f(X_i, \beta) + \epsilon_i$

Matrix decomposition of *TD* matrix

For bag-of-words representation X w. documents in columns and words in rows, such that each entry X_{ij} is the ith word in the jth column, we solve

$$\underset{W,H}{\operatorname{argmin}} \parallel X - WH \parallel_F^2 \quad s.t.\ W, H \geq 0$$

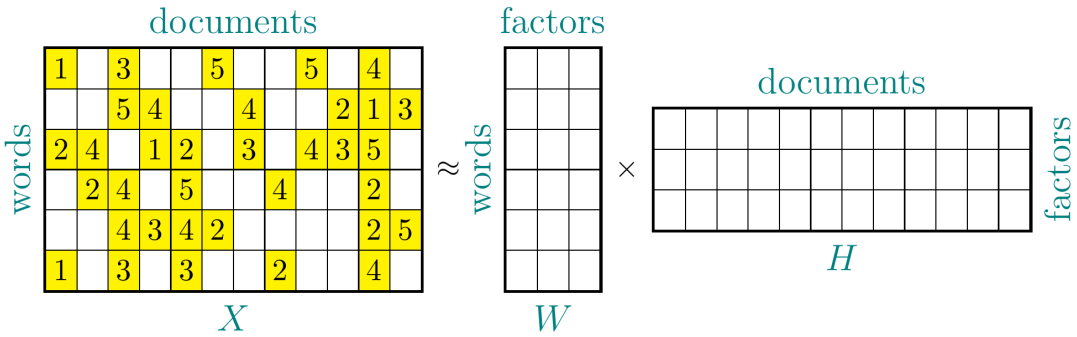

Non-negative matrix factorization of *TD* matrix X

WH is a low-rank approximation of the data (++fewer factors than documents and words) and each document is the weighted sum of columns in W, "document factors" $\sim$ themes or topics

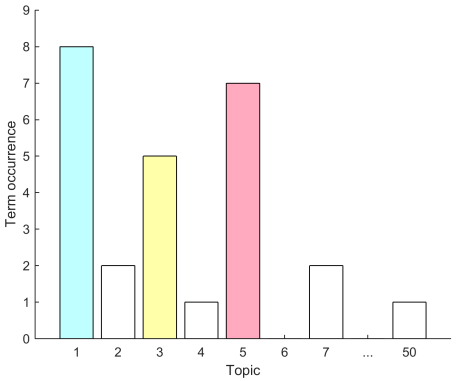

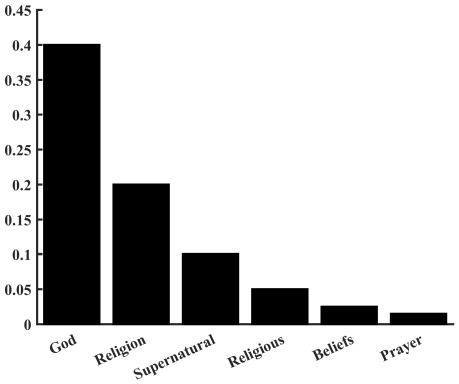

Document representation from NMF

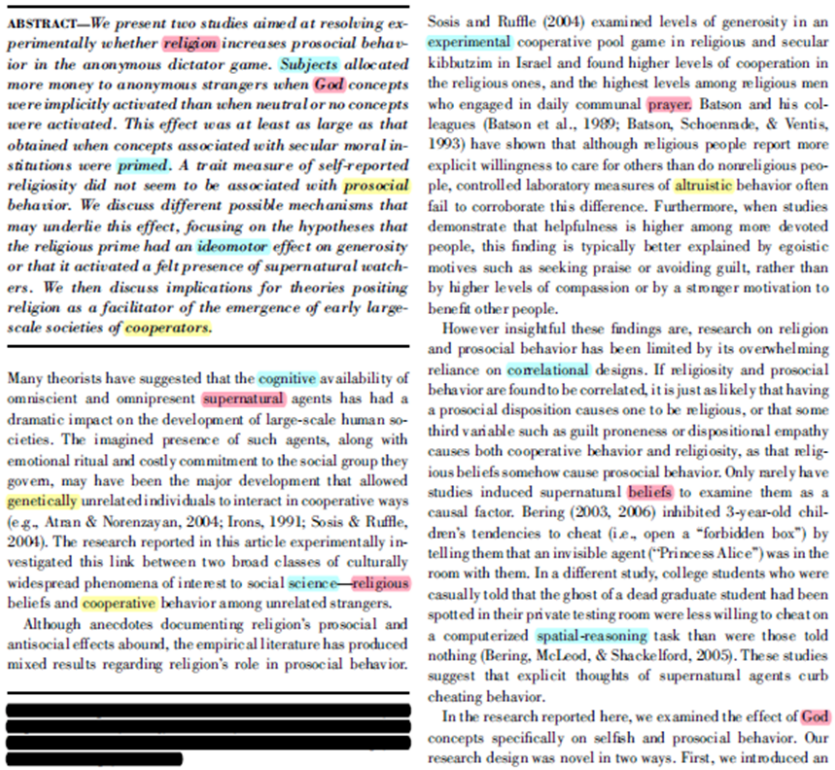

ABSTRACT—*We present two studies aimed at resolving experimentally whether religion increases prosocial behavior in the anonymous dictator game. Subjects allocated more money to anonymous strangers when God concepts were implicitly activated than when neutral or no concepts were activated. This effect was at least as large as that obtained when concepts associated with secular moral institutions were primed. A trait measure of self-reported religiosity did not seem to be associated with prosocial behavior. We discuss different possible mechanisms that may underlie this effect, focusing on the hypotheses that the religious prime had an ideomotor effect on generosity or that it activated a felt presence of supernatural watchers. We then discuss implications for theories positing religion as a facilitator of the emergence of early large-scale societies of cooperators.*

Many theorists have suggested that the cognitive availability of omniscient and omnipresent supernatural agents has had a dramatic impact on the development of large-scale human societies. The imagined presence of such agents, along with emotional ritual and costly commitment to the social group they govern, may have been the major development that allowed genetically unrelated individuals to interact in cooperative ways (e.g., Atran & Norenzayan, 2004; Irons, 1991; Sosis & Ruffle, 2004). The research reported in this article experimentally investigated this link between two broad classes of culturally widespread phenomena of interest to social science—religious beliefs and cooperative behavior among unrelated strangers.

Although anecdotes documenting religion's prosocial and antisocial effects abound, the empirical literature has produced mixed results regarding religion's role in prosocial behavior.

[redacted]

Sosis and Ruffle (2004) examined levels of generosity in an experimental cooperative pool game in religious and secular kibbutzim in Israel and found higher levels of cooperation in the religious ones, and the highest levels among religious men who engaged in daily communal prayer. Batson and his colleagues (Batson et al., 1989; Batson, Schoenrade, & Ventis, 1993) have shown that although religious people report more explicit willingness to care for others than do nonreligious people, controlled laboratory measures of altruistic behavior often fail to corroborate this difference. Furthermore, when studies demonstrate that helpfulness is higher among more devoted people, this finding is typically better explained by egoistic motives such as seeking praise or avoiding guilt, rather than by higher levels of compassion or by a stronger motivation to benefit other people.

However insightful these findings are, research on religion and prosocial behavior has been limited by its overwhelming reliance on correlational designs. If religiosity and prosocial behavior are found to be correlated, it is just as likely that having a prosocial disposition causes one to be religious, or that some third variable such as guilt proneness or dispositional empathy causes both cooperative behavior and religiosity, as that religious beliefs somehow cause prosocial behavior. Only rarely have studies induced supernatural beliefs to examine them as a causal factor. Bering (2003, 2006) inhibited 3-year-old children's tendencies to cheat (i.e., open a "forbidden box") by telling them that an invisible agent ("Princess Alice") was in the room with them. In a different study, college students who were casually told that the ghost of a dead graduate student had been spotted in their private testing room were less willing to cheat on a computerized spatial-reasoning task than were those told nothing (Bering, McLeod, & Shackelford, 2005). These studies suggest that explicit thoughts of supernatural agents curb cheating behavior.

In the research reported here, we examined the effect of God concepts specifically on selfish and prosocial behavior. Our research design was novel in two ways. First, we introduced an

Original document

You can use the 'topics' as predictors $x_1, x_2, ...x_n$ in $Y_i = f(X_i, \beta) + \epsilon_i$ where $n < m$ for word-based regression

Text Classification

- Given labeled data, a classification algorithm will output a solution that categorizes new examples → associate labels with subsets of the data (c.f., logistic regression)
- data (features) with class values (∼ labeled data), excellent opportunity to make use of metadata (e.g., reviews scores, evaluations &c)

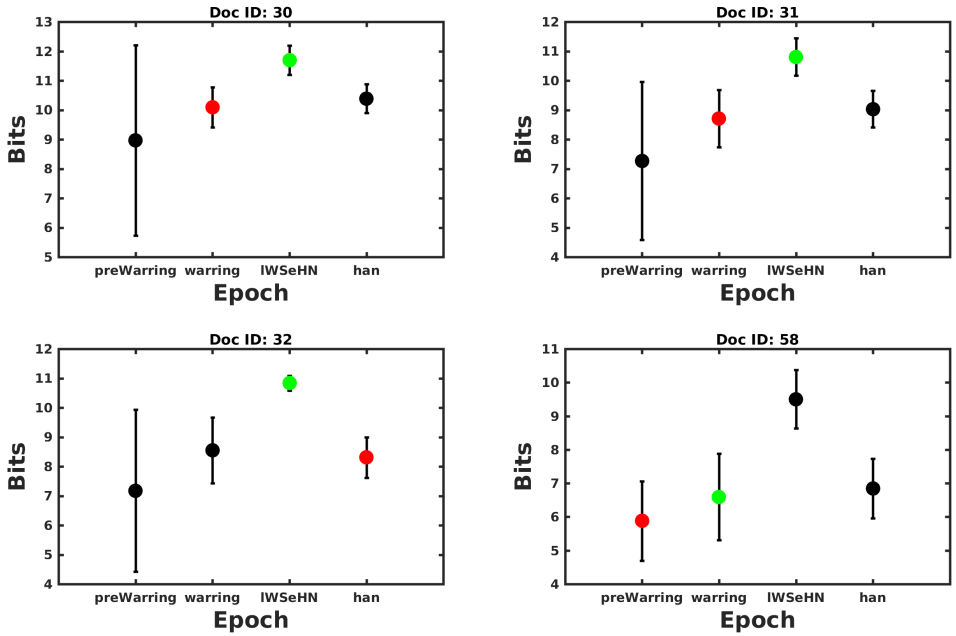

Source classification of Chinese historical documents

Kristoffer L Nielbo
chcaa.au.dk
chcaa.io

Naive Bayes

A simple and very popular probability learning model that can be implemented very efficiently

The probability of a document d being in class c, $P(c \mid d)$ is computed as:

$$P(c \mid d) \propto P(c) \prod_{i=1}^{m} P(t_i \mid c) \tag{1}$$

and the class of a document d is then computed as:

$$c_{MAP} = arg\ max_{c \in \{c_1, c_2\}} P(c \mid d) \tag{2}$$

Naive assumption that the presence/absence of a feature is completely independent of other features.

Exploratory analysis with Deep Neural Nets

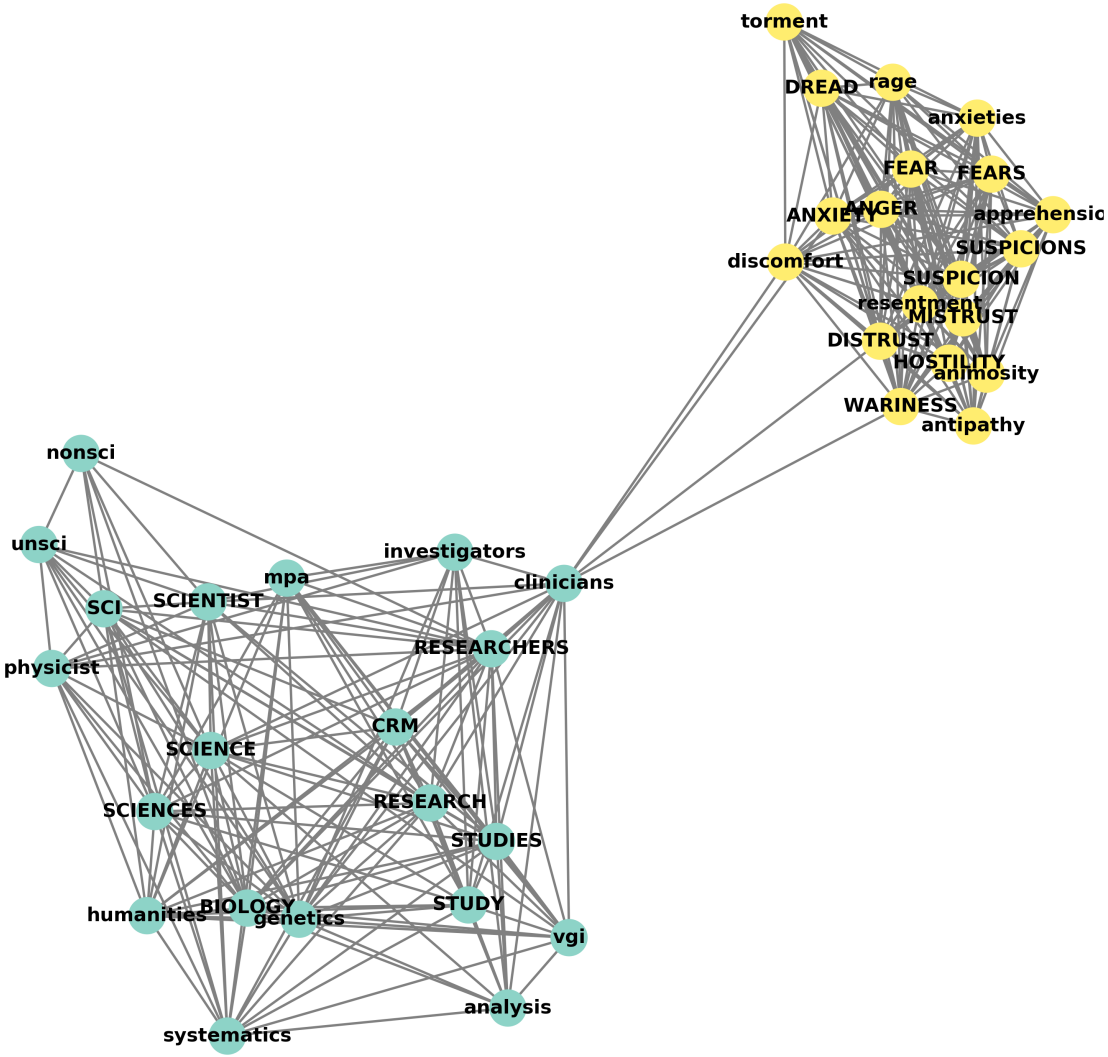

Sparse query graph for MISTRUST and SCIENCE

Distributed word representation

For a discrete input $w_0, w_1, ..., w_n$, we train a simple feedforward network such that

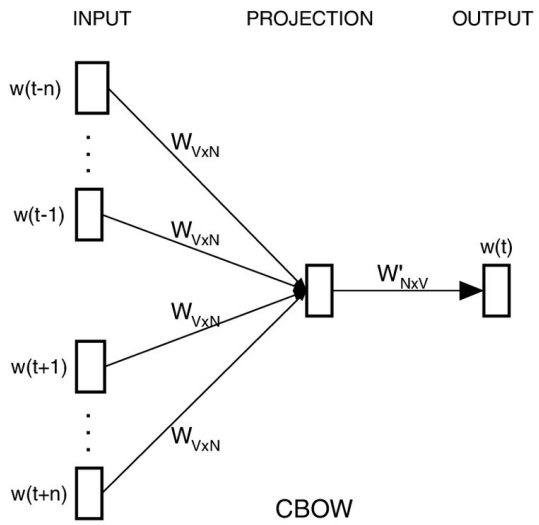

Figure 1: CBOW architecture for learning neural embeddings at the word level

$$\frac{1}{T}\sum_{t=1}^{T}\sum_{-n\leq j\leq n}\log p(w_t|w_{t+j}), \quad j\neq 0$$

Semantic similarity between any two word embeddings, A and B, can then be measured as their angular similarity

$$1-\frac{\cos^{-1}\left(\frac{A\cdot B}{\|A\|\,\|B\|}\right)}{\pi}$$

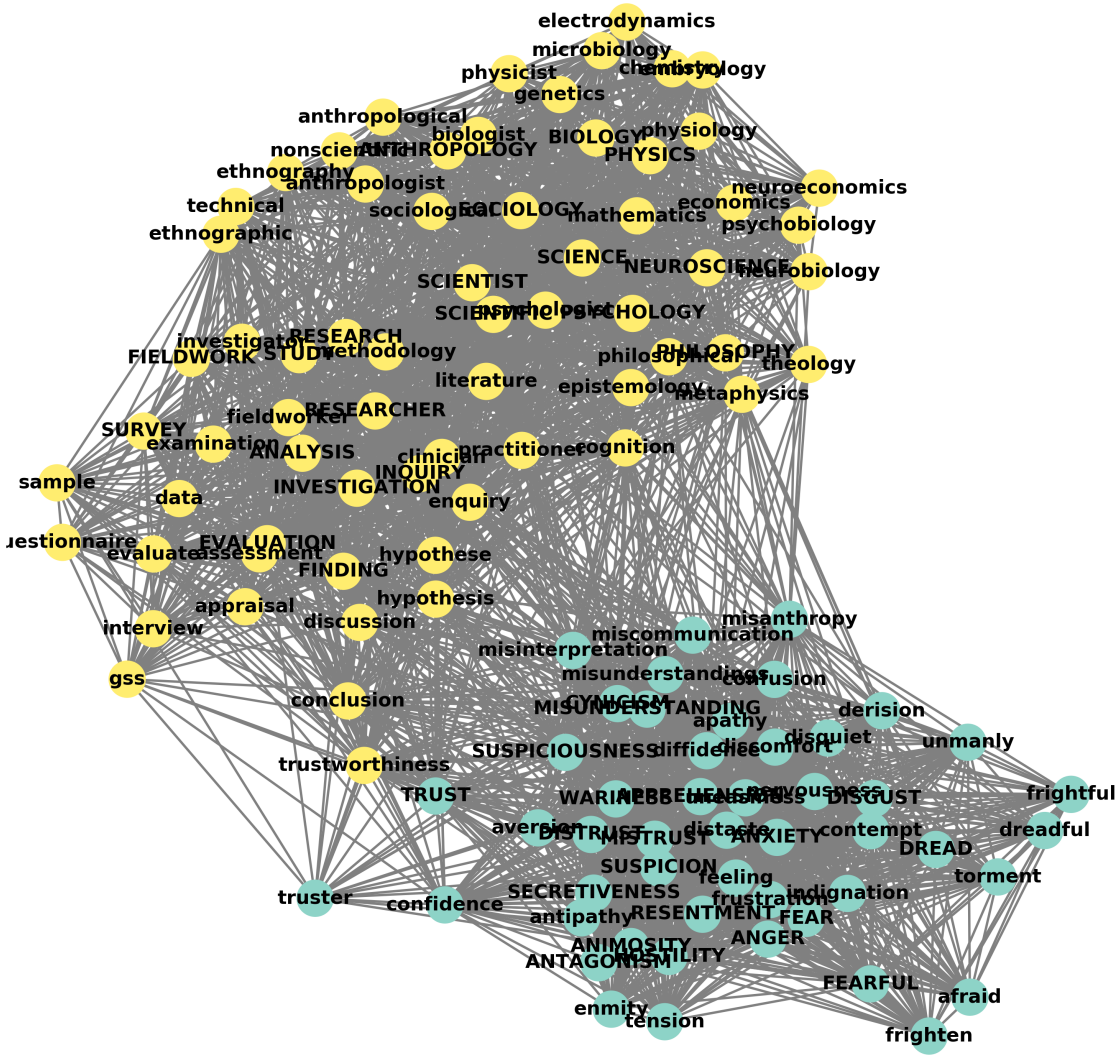

Dense query graph for MISTRUST and SCIENCE

Thank you for your attention

chcaa@cas.au.dk
chcaa.io

slides: http://knielbo.github.io/files/kln_mginno.pdf